

# AIM-7 Performance Automation

(available in PDF format as part of the install package)

**Dec 13, 2012**  
**Scott Norton**



# Prior to installing AIM-7 – *if needed*



```
rpm --import /etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

```
# Update system if desired
```

```
cd /tmp
```

```
wget http://satserv.usa.hp.com/pub/bootstrap/bootstrap.sh
```

```
sh bootstrap.sh
```

```
# packages required to install AIM-7
```

```
yum -y install libaio-devel.x86_64
```

**# required**

```
yum -y install automake-1.11.1-1.2.el6.noarch
```

**# required**

```
yum -y install make
```

**# required**

```
yum -y install libtool-2.2.6-15.5.el6.x86_64
```

**# required**

```
# Some useful tools
```

```
yum -y install numactl
```

```
yum -y install powertop
```

```
yum -y install perf
```

# Installing AIM-7



- Remove old `reaim` directory contents (or save them to another directory):
  - `rm -rf /usr/local/share/reaim/*`
- Obtain the latest AIM-7 revision from:
  - <http://mint-autotest.cce.hp.com/kernels/perf/>
- Place it on your test machine in:
  - `/usr/local/share/reaim/.`
- Install
  - `cd /usr/local/share/reaim`
  - `tar xvjf hp-osdl-aim-7.0.1.13.hp-rev7.tar.gz`
  - `./osdl-aim-7/setup`
- By default it will setup ram-based filesystems for use in `/t0, /t1, ..., /t15`
  - If you want different mount points follow the instructions displayed at the end of the setup output.
- ‘`reaim`’ and the new ‘`run_reaim`’ script will be in `/usr/local/share/reaim`
  - They can be run from anywhere – you don’t have to be in this directory

# Kernel Boot Parameters



- Add the following to the boot options to your kernel(s) in `/etc/grub.conf`:
  - `ramdisk_size=131072`
- This value matches the ram-based filesystem sizes that are specified in the file:
  - `/usr/local/share/ream/osdl-aim-7/Support_scripts/master_disk_list`
- The default `ramdisk_size` in Linux, if not specified, is 16384 (in K bytes) = 16M
- If you don't do this you will see several "WARNING:" messages when the automated scripts create ram-based filesystems
  - With the current defaults in the automation scripts the AIM-7 workloads seem to be able to run fine without adding `ramdisk_size=131072` to `/etc/grub.conf` (but you will get a bunch of warnings)

# AIM-7 Workloads



- AIM-7 has 14 workloads that we use:
  - all\_untime
  - alltests
  - compute
  - custom
  - dbase
  - disk (*can't run with ramfs*)
  - five\_sec
  - fserver
  - high\_systime
  - long
  - new\_dbase
  - new\_fserver
  - shared
  - short
- Workloads are specified in files located in `/usr/local/share/ream` such as:
  - `workfile.all_untime`
  - `workfile.dbase`
  - `workfile.fserver`

# Running AIM-7: The Traditional Way



- Use the `ream` command. For example:
  - `./ream -s100 -e1000 -t -j100 -i100 -y -f ./workfile.fserver`
  - `./ream -s100 -e1000 -t -j100 -i100 -y -f ./workfile.dbase`
- Options:
  - `-s<num>` is the starting number of users
  - `-e<num>` is the ending number of users
  - `-i<num>` is the increment in the number of users per run
  - `-j<num>` is jobs per user (which is 100 by default)
  - `-t` disables adaptive timer so that `-s` to `-e` will always be in `-i` increments
  - `-y` do not execute "sync"
- Results are output to stdout as well as to the following files:
  - `multiuser.ss` space separated "human readable" results
  - `ream.csv` comma separated results to put into an Excel spreadsheet

# Running AIM-7: The Traditional Way



```
[root@vlab380j reaim]# ./ream -s100 -e1000 -t -j100 -i100 -y -f ./workfile.dbase
```

```
Using default config file 'ream.config'  
No logfile prefix specified, using default 'ream'
```

```
Not executing sync during disk activity---faster execution.
```

```
DEBUG: Number of directories is 16
```

```
REAM Workload
```

```
Times are in seconds - Child times from tms.cstime and tms.cutime
```

Num	Parent	Child	Child	Jobs per	Jobs/min/	Std_dev	Std_dev	JTI
Forked	Time	SysTime	UTime	Minute	Child	Time	Percent	
100	4.86	2.04	56.11	122222.22	1222.22	0.62	13.93	86
200	9.75	3.91	112.68	121846.15	609.23	1.93	21.28	78
300	14.64	6.03	168.87	121721.31	405.74	0.93	6.62	93
400	19.51	8.07	225.31	121783.70	304.46	0.52	2.70	97
500	24.37	9.88	281.70	121871.15	243.74	0.28	1.17	98
600	29.18	11.66	337.69	122138.45	203.56	0.64	2.21	97
700	34.09	13.61	394.74	121971.25	174.24	0.35	1.04	98
800	39.00	15.79	451.29	121846.15	152.31	0.45	1.17	98
900	43.81	17.61	507.12	122026.93	135.59	0.45	1.04	98
1000	48.63	19.71	562.61	122146.82	122.15	0.50	1.05	98

```
Max Jobs per Minute 122222.2
```

# Running AIM-7: The Traditional Way

*ram-based FS vs. disk-based FS*



- We typically run AIM-7 with a ram-based file system (unless we really want to test a specific disk device).
- To use a ram-based file system follow the example below:

```
./osdl-aim-7/Support_scripts/do_make_file_systems
```

```
./osdl-aim-7/Support_scripts/do_mount_file_systems
```

```
./osdl-aim-7/Support_scripts/do_clean_file_systems # remove existing files
```

```
./reaim -s100 -e1000 -j100 -i100 -y -f ./workfile.fserver
```

```
./osdl-aim-7/Support_scripts/do_clean_file_systems # remove existing files
```

```
./osdl-aim-7/Support_scripts/do_umount_file_systems
```

- Note: files persist on a file system such as `ext4` with `ramdisk` between `umount` and `(re)mount`. This is not true for `ramfs`.
- Note: `ramfs` releases memory back to the system for reuse after `umount`, `ext4` with `ramdisk` does not



# Running AIM-7: The Traditional Way



- If you are analyzing the performance of a specific workload at a specific user level load, the “Traditional Way” works just fine
- However, if you want to run
  - All or many workloads
  - Use ram-based file systems
  - Have the full range of data points from 10 users to 2000 users
  - Have averaged results instead of single data points

and then have to save away the results and do this for multiple kernels to compare can be a lot of manual work

# Running AIM-7: `run_reaim` script



- The new `run_reaim` script will automate all of this. By default it:
  - Runs all workloads (except aim9)
  - Runs 10 to 90 users in increments of 10 and 100 to 2000 users in increments of 100
  - Uses a ram-based file system
  - Saves data in `<workload>.ssv` and `<workload>.csv` files in the `./results` directory (for example, `./results/compute.ssv` and `./results/compute.csv`)
  - Successive `run_reaim` runs add data to the end of the `.ssv` and `.csv` files
- Output from the `run_reaim` script (including the `reaim` command output) is redirected to `./results/run_reaim.out`
  - This file is also copied to `./results/run_reaim.out.$$` so future runs won't overwrite your previous output.
- Look through the `./results/run_reaim.out` file for any errors after running the workloads

# Running ALM-7: `run_reaim` script



- After running `run_reaim` the output directory `./results` looks like:

<code>all_untime.csv</code>	<code>custom.ssv</code>	<code>fserver.csv</code>	<code>new_dbase.ssv</code>	<code>short.csv</code>
<code>all_untime.ssv</code>	<code>dbase.csv</code>	<code>fserver.ssv</code>	<code>new_fserver.csv</code>	<code>short.ssv</code>
<code>alltests.csv</code>	<code>dbase.ssv</code>	<code>high_systemtime.csv</code>	<code>new_fserver.ssv</code>	
<code>alltests.ssv</code>	<code>disk.csv</code>	<code>high_systemtime.ssv</code>	<code>run_reaim.out</code>	
<code>compute.csv</code>	<code>disk.ssv</code>	<code>long.csv</code>	<code>run_reaim.out.3868</code>	
<code>compute.ssv</code>	<code>five_sec.csv</code>	<code>long.ssv</code>	<code>shared.csv</code>	
<code>custom.csv</code>	<code>five_sec.ssv</code>	<code>new_dbase.csv</code>	<code>shared.ssv</code>	

- The `.ssv` files are space separated values that are “human readable”
  - The `.csv` files are comma separated values that can be read into an Excel spreadsheet
    - More on this later in this slide set
- Each time `run_reaim` is run it will add data to the end of the `.csv` and `.ssv` files
  - When you want to start fresh, either remove all these files or specify the `-r` option to `run_reaim`

# Running AIM-7: run\_reaim script

## Sample dbase.ssv output file



Title: hydrazine0 DL980 80-core  
Kernel: hydrazine0 RHEL 2.6.32-279.el6.x86\_64 HT-disabled ramfs  
Date: Thu Nov 8 02:06:16 MST 2012  
Command: /usr/local/share/reaim/osdl-aim-7/reaim -s10 -e2000 -t -j100 -i[10|100] -y -f ./workfile.dbase  
Workload: dbase  
Time: 2 mins, 57 secs

Forks	Jobs/min	Jobs/min/ child	Jobs/sec/ child	Time: parent	Time: childU	Time: childS	Time: std_dev	JTI	Running child time	
									:max	:min
10	99000.00	9900.00	165.00	0.60	4.28	0.22	0.07	87.00	0.60	0.44
20	237600.00	11880.00	198.00	0.50	8.65	0.48	0.02	96.00	0.50	0.44
30	187578.95	6252.63	104.21	0.95	13.14	0.64	0.14	78.00	0.95	0.44
40	440000.00	11000.00	183.33	0.54	17.74	0.87	0.02	94.00	0.53	0.44
50	550000.00	11000.00	183.33	0.54	22.67	1.11	0.02	95.00	0.54	0.46
60	548307.69	9138.46	152.31	0.65	27.71	1.31	0.05	91.00	0.65	0.49
70	540000.00	7714.29	128.57	0.77	32.40	1.59	0.06	91.00	0.77	0.46
80	505531.91	6319.15	105.32	0.94	36.93	1.74	0.11	87.00	0.94	0.46
90	712800.00	7920.00	132.00	0.75	41.90	2.04	0.07	89.00	0.75	0.51
100	813698.63	8136.99	135.62	0.73	46.86	2.33	0.03	95.00	0.73	0.60
200	886567.16	4432.84	73.88	1.34	93.68	4.73	0.04	96.00	1.34	1.10
300	909183.67	3030.61	50.51	1.96	140.62	7.20	0.06	96.00	1.96	1.65
400	876752.77	2191.88	36.53	2.71	187.50	9.35	0.08	96.00	2.71	2.16
500	931034.48	1862.07	31.03	3.19	234.80	11.71	0.11	96.00	3.19	2.53
600	909183.67	1515.31	25.26	3.92	281.97	13.99	0.15	95.00	3.92	2.70
700	932287.00	1331.84	22.20	4.46	329.00	16.34	0.14	96.00	4.45	3.73
800	939130.43	1173.91	19.57	5.06	376.05	18.97	0.19	96.00	5.06	3.78
900	936252.19	1040.28	17.34	5.71	423.26	21.19	0.19	96.00	5.71	4.73
1000	913846.15	913.85	15.23	6.50	470.60	23.67	0.23	96.00	6.49	5.15
1100	860869.57	782.61	13.04	7.59	518.45	25.68	0.28	96.00	7.59	6.08
1200	897733.00	748.11	12.47	7.94	566.01	28.19	0.29	96.00	7.93	6.35
1300	937135.92	720.87	12.01	8.24	612.61	31.01	0.30	96.00	8.24	5.77
1400	924000.00	660.00	11.00	9.00	660.13	33.43	0.36	95.00	9.00	6.66
1500	936908.52	624.61	10.41	9.51	707.30	35.70	0.38	95.00	9.51	7.63
1600	939130.43	586.96	9.78	10.12	754.62	38.29	0.35	96.00	10.11	8.33
1700	942857.14	554.62	9.24	10.71	801.78	40.33	0.42	95.00	10.70	8.36
1800	936252.19	520.14	8.67	11.42	849.15	42.93	0.46	95.00	11.42	8.72
1900	938154.61	493.77	8.23	12.03	897.33	45.33	0.43	96.00	12.03	9.36
2000	939873.42	469.94	7.83	12.64	944.26	47.93	0.50	95.00	12.64	8.05

Highest JPM: 942857.14 at 1700 users

\*\*\*\*\* Data Separator \*\*\*\*\*

# Running AIM-7: run reaim script

## *Sample dbase.csv output file – to put in an excel spreadsheet*



```
Title:,hydrazine0 DL980 80-core w/ ramfs
Kernel:,hydrazine0 RHEL 2.6.32-279.el6.x86_64 HT-disabled ramfs
Date:.,Thu Nov 8 02:06:16 MST 2012
Command:., /usr/local/share/ream/osdl-aim-7/ream -s10 -e2000 -j100 -t -i[10|100] -y -f ./workfile.dbase
Workload:.,dbase
Time:.,mins:.,2,secs:.,57
Forks,JPM,JPM_C,JPS_C,parent_tm,childU_tm,childS_tm,std_dev,JTI,max_c,min_c
10,99000.00,9900.00,165.00,0.60,4.28,0.22,0.07,87.00,0.60,0.44
20,237600.00,11880.00,198.00,0.50,8.65,0.48,0.02,96.00,0.50,0.44
30,187578.95,6252.63,104.21,0.95,13.14,0.64,0.14,78.00,0.95,0.44
40,440000.00,11000.00,183.33,0.54,17.74,0.87,0.02,94.00,0.53,0.44
50,550000.00,11000.00,183.33,0.54,22.67,1.11,0.02,95.00,0.54,0.46
60,548307.69,9138.46,152.31,0.65,27.71,1.31,0.05,91.00,0.65,0.49
70,540000.00,7714.29,128.57,0.77,32.40,1.59,0.06,91.00,0.77,0.46
80,505531.91,6319.15,105.32,0.94,36.93,1.74,0.11,87.00,0.94,0.46
90,712800.00,7920.00,132.00,0.75,41.90,2.04,0.07,89.00,0.75,0.51
100,813698.63,8136.99,135.62,0.73,46.86,2.33,0.03,95.00,0.73,0.60
200,886567.16,4432.84,73.88,1.34,93.68,4.73,0.04,96.00,1.34,1.10
300,909183.67,3030.61,50.51,1.96,140.62,7.20,0.06,96.00,1.96,1.65
400,876752.77,2191.88,36.53,2.71,187.50,9.35,0.08,96.00,2.71,2.16
500,931034.48,1862.07,31.03,3.19,234.80,11.71,0.11,96.00,3.19,2.53
600,909183.67,1515.31,25.26,3.92,281.97,13.99,0.15,95.00,3.92,2.70
700,932287.00,1331.84,22.20,4.46,329.00,16.34,0.14,96.00,4.45,3.73
800,939130.43,1173.91,19.57,5.06,376.05,18.97,0.19,96.00,5.06,3.78
900,936252.19,1040.28,17.34,5.71,423.26,21.19,0.19,96.00,5.71,4.73
1000,913846.15,913.85,15.23,6.50,470.60,23.67,0.23,96.00,6.49,5.15
1100,860869.57,782.61,13.04,7.59,518.45,25.68,0.28,96.00,7.59,6.08
1200,897733.00,748.11,12.47,7.94,566.01,28.19,0.29,96.00,7.93,6.35
1300,937135.92,720.87,12.01,8.24,612.61,31.01,0.30,96.00,8.24,5.77
1400,924000.00,660.00,11.00,9.00,660.13,33.43,0.36,95.00,9.00,6.66
1500,936908.52,624.61,10.41,9.51,707.30,35.70,0.38,95.00,9.51,7.63
1600,939130.43,586.96,9.78,10.12,754.62,38.29,0.35,96.00,10.11,8.33
1700,942857.14,554.62,9.24,10.71,801.78,40.33,0.42,95.00,10.70,8.36
1800,936252.19,520.14,8.67,11.42,849.15,42.93,0.46,95.00,11.42,8.72
1900,938154.61,493.77,8.23,12.03,897.33,45.33,0.43,96.00,12.03,9.36
2000,939873.42,469.94,7.83,12.64,944.26,47.93,0.50,95.00,12.64,8.05
High JPM:., 942857.14,1700,users,
***** , ***** Data Separator *****
```

# Running AIM-7: run\_reaim options



Usage: run\_reaim

```
-k <kernel>      # name to appear on graph line
-t <title>       # name to appear on graph title (i.e., system name/type)
-d <log_dir>     # output logging directory (default: "./results")
-c <cmd>         # reaim preface (for ex: -c "chrt --fifo 3")
-r             # remove old logging directory files first
-h             # help menu

-ramfs          # use a ramfs for workloads
-ext4_ramdisk   # use an ext4 ramdisk for workloads (default)
-ext3_ramdisk   # use an ext3 ramdisk for workloads
-ext2_ramdisk   # use an ext2 ramdisk for workloads
-diskfs         # use a diskfs for workloads

-avg <num>      # average <num> runs

-perf           # use perf record/report - output is in ./results/<wkload>.perf

-x1            # do *not* run 10-90      users option
-x2            # do *not* run 100-1000   users option
-x3            # do *not* run 1100-2000  users option
-x4            # do *not* run $OPTS4     users option (currently unused)
-x5            # do *not* run $OPTS5     users option (currently unused)
-x6            # do *not* run $OPTS6     users option (currently unused)

-<workload>     # do *not* run <workload>

+<workload>     # only run <workload>

-aim9only      # only run aim9
```

-k <kernel> and -t <title> are required

Avail Workloads:	short	dbase	new_dbase	compute	all_utime
	disk	five_sec	long	shared	alltests
	custom	new_fserver	fserver	high_sysime	aim9

# Running AIM-7: run\_reaim examples



- Run with default options (all workloads, 10-2000 users, ext4 ram-based fs):  
`./run_reaim -k "my kernel" -t "my machine"`
- Run all workloads with only 100 to 1000 users:  
`./run_reaim -k "my kernel" -t "my machine" -x1 -x3`
- Run only the dbase and fserver workloads:  
`./run_reaim -k "my kernel" -t "my machine" +dbase +fserver`
- Run all workloads except short and custom:  
`./run_reaim -k "my kernel" -t "my machine" -short -custom`
- Run only the dbase workload with only 100 to 1000 users:  
`./run_reaim -k "my kernel" -t "my machine" -x1 -x2 +dbase`
- Run all workloads with disk-based filesystems (as opposed to a ram-based fs)  
`./run_reaim -k "my kernel" -t "my machine" -diskfs`
- Run all workloads using real-time scheduling:  
`./run_reaim -k "my kernel" -t "my machine" -c "chrt --fifo 3"`

# Running AIM-7: `run_reaim` examples



- At this point in the presentation show a live example of installing AIM-7 and using the `run_reaim` script...

```
- cd /usr/local/share/reaim
- tar xvjf hp-osdl-aim-7.0.1.13.hp-rev7.tar.gz
- ./osdl-aim-7/setup
- ./run_reaim -k "my kernel" -t "my machine" +dbase
- more ./results/run_reaim.out
- more ./results/dbase.ssv
- more ./results/dbase.csv
```



# Running AIM-7: `run_reaim` script

*Average option*



- To run a workload multiple times and average each data point use the “`-avg <num>`” option
- Each individual run is still added to the end of `./results/<workload>.ssv` and `./results/<workload>.csv`
- Two new files are created that contain the averaged data from `<num>` runs:
  - `./results/<workload>.avg.ssv`
  - `./results/<workload>.avg.csv`
- As with the other results files, future “`-avg <num>`” runs add to the end of the existing `<workload>.avg.ssv` and `<workload>.avg.csv` files

# Running AIM-7: `run_reaim -avg <num>`

## Sample `dbase.avg.ssv` output file



Title: hydrazine0 DL980 80-core w/ ramfs  
Kernel: average - hydrazine0 RHEL 2.6.32-279.el6.x86\_64 HT-disabled ramfs  
Date: Sat Nov 17 18:27:56 MST 2012  
Command: ./reaim -s10 -e2000 -t -j100 -i[10|100] -y -f ./workfile.dbase  
Workload: dbase  
Time:

Jobs	Avg	Median	High	Low	H/L_Diff	H/L_Pct	Val1	Diff1	Pct1	Val2	Diff2	Pct2	...	...
----	----	-----	----	----	-----	-----	----	-----	-----	----	-----	-----	----	----
10	113464.18	126382.98	132000	63870	68130	51.61	63870.97	-49593.21	-43.70%	72439.02	-41025.16	-36.15%	----	----
20	198051.83	235270.59	252765	113142	139623	55.23	113142.86	-84908.97	-42.87%	127741.94	-70309.89	-35.50%		
30	283815.23	330453.30	356400	147272	209128	58.67	147272.73	-136542.50	-48.10%	173009.71	-110805.52	-39.04%		
40	388159.05	420827.58	484897	220000	264897	54.62	220000.00	-168159.05	-43.32%	240000.00	-148159.05	-38.16%		
50	461464.75	480156.45	582352	270000	312352	53.63	270000.00	-191464.75	-41.49%	309375.00	-152089.75	-32.95%		
60	571000.90	604067.80	698823	307241	391582	56.03	307241.38	-263759.52	-46.19%	318214.29	-252786.61	-44.27%		
70	664880.87	784528.30	799615	270000	529615	66.23	270000.00	-394880.87	-59.39%	374594.59	-290286.28	-43.65%		
80	753332.80	800067.91	896603	409655	486948	54.31	409655.17	-343677.63	-45.62%	552558.14	-200774.66	-26.65%		
90	787283.25	810000.00	848571	540000	308571	36.36	540000.00	-247283.25	-31.40%	763714.29	-23568.96	-2.99%		
100	778201.47	813698.63	836619	456923	379696	45.38	456923.08	-321278.39	-41.28%	751898.73	-26302.74	-3.37%		
200	821775.11	845562.31	893233	686705	206528	23.12	686705.20	-135069.91	-16.43%	694736.84	-127038.27	-15.45%		
300	854743.19	886567.16	904568	690697	213871	23.64	690697.67	-164045.52	-19.19%	778165.94	-76577.25	-8.95%		
400	900132.93	906870.23	917374	854676	62698	6.83	854676.26	-45456.67	-5.04%	864000.00	-36132.93	-4.01%		
500	916026.63	915256.41	928125	900000	28125	3.03	900000.00	-16026.63	-1.74%	905487.80	-10538.83	-1.15%		
600	919634.54	918562.80	935433	904568	30865	3.29	904568.53	-15066.01	-1.63%	911508.95	-8125.59	-0.88%		
700	927521.82	932287.00	942857	888461	54396	5.76	888461.54	-39060.28	-4.21%	905882.35	-21639.47	-2.33%		
800	925719.50	932679.99	939130	842553	96577	10.28	842553.19	-83166.31	-8.98%	924513.62	-1205.88	-0.13%		
900	930549.78	933799.83	942857	910732	32125	3.40	910732.54	-19817.24	-2.12%	912286.69	-18263.09	-1.96%		
1000	933197.42	935435.39	942857	912442	30415	3.22	912442.40	-20755.02	-2.22%	919504.64	-13692.78	-1.46%		
1100	941737.97	942857.14	945586	933428	12158	1.28	933428.57	-8309.40	-0.88%	936103.15	-5634.82	-0.59%		
1200	937213.62	939130.43	946613	923316	23297	2.46	923316.06	-13897.56	-1.48%	928125.00	-9088.62	-0.96%		
1300	939440.62	941133.80	944009	925899	18110	1.91	925899.28	-13541.24	-1.44%	934866.83	-4573.79	-0.48%		
1400	936937.54	939662.21	945000	907860	37140	3.93	907860.26	-29077.38	-3.10%	933333.33	-3604.21	-0.38%		
1500	940625.11	940865.89	945859	936908	8951	0.94	936908.52	-3716.59	-0.39%	936908.52	-3716.59	-0.39%		
1600	937469.32	939131.35	943793	924513	19280	2.04	924513.62	-12955.70	-1.38%	928125.00	-9344.32	-0.99%		
1700	937144.14	941099.72	942857	920510	22347	2.37	920510.48	-16633.66	-1.77%	929834.25	-7309.89	-0.78%		
1800	936369.48	939956.22	942026	903040	38986	4.13	903040.54	-33328.94	-3.55%	934615.38	-1754.10	-0.18%		
1900	939005.55	938935.11	942070	934271	7799	0.82	934271.52	-4734.03	-0.50%	935820.90	-3184.65	-0.33%		
2000	940312.81	940617.58	944356	936908	7448	0.78	936908.52	-3404.29	-0.36%	937647.99	-2664.82	-0.28%		

Highest JPM Average: 941737 at 1100 users  
Highest JPM Median: 942857 at 1100 users  
Highest JPM Single: 946613 at 1200 users

\*\*\*\*\* Data Separator \*\*\*\*\*

# Running AIM-7: `-perf` option

- The `-perf` option causes each workload to be run with:
  - `perf record -a -g -s`
- At the end of each workload run a perf report is generated with:
  - `perf report -n --stdio`
- Output files are located in the directory `./results/<workload>.perf`
- Three perf output files per workload run:
  - `<workload>.10_90.perf.<pid>.out` # 10-90 users
  - `<workload>.100_1000.perf.<pid>.out` # 100-1000 users
  - `<workload>.1100_2000.perf.<pid>.out` # 1100-2000 users
- The `<pid>` is added to the filename so future runs will not overwrite previous data

# Sysctl values



- Running the AIM-7 workloads requires some sysctl values to be increased depending on what your system currently has set
- The ones that we have found so far have been incorporated into `run_reaim` so you don't have to do it:
  - `kernel.sem`
  - `kernel.shmmax`
  - `kernel.shmall`
  - `net.netfilter.nf_conntrack_max`
  - `net.ipv4.tcp_fin_timeout`
  - `net.ipv4.tcp_keepalive_time`
  - `net.ipv4.tcp_keepalive_intvl`
  - `net.ipv4.tcp_tw_recycle`
  - `net.ipv4.tcp_tw_reuse`
- If `run_reaim` modifies any of these values it will be displayed at the beginning of the `./results/run_reaim.out` output file
- If you find others that need to be added please let us know

# Different `reaim` run options



- If you want different run options to `reaim`, you can make a copy of `run_reaim` and modify the new copy.
- In your copy look for the following lines near the top of the file:

```
OPTS1="-s10    -e90    -t -j100 -i10    -y"  
OPTS2="-s100   -e1000  -t -j100 -i100   -y"  
OPTS3="-s1100  -e2000  -t -j100 -i100   -y"  
  
# not defined now - user can modify as desired  
OPTS4=""  
OPTS5=""  
OPTS6=""
```

- You can modify the above defines for your specific options (without changing the existing behavior). As an example, you can do something like:

```
OPTS4="-s2100  -e3000  -t -j100 -i100   -y"
```

to run 2100 to 3000 users since the current values stop at 2000 users


- Once modified you can run your script
  - Running with options `-x1 -x2 -x3` will cause it to run only with your new options

# Graphing AIM-7 Output

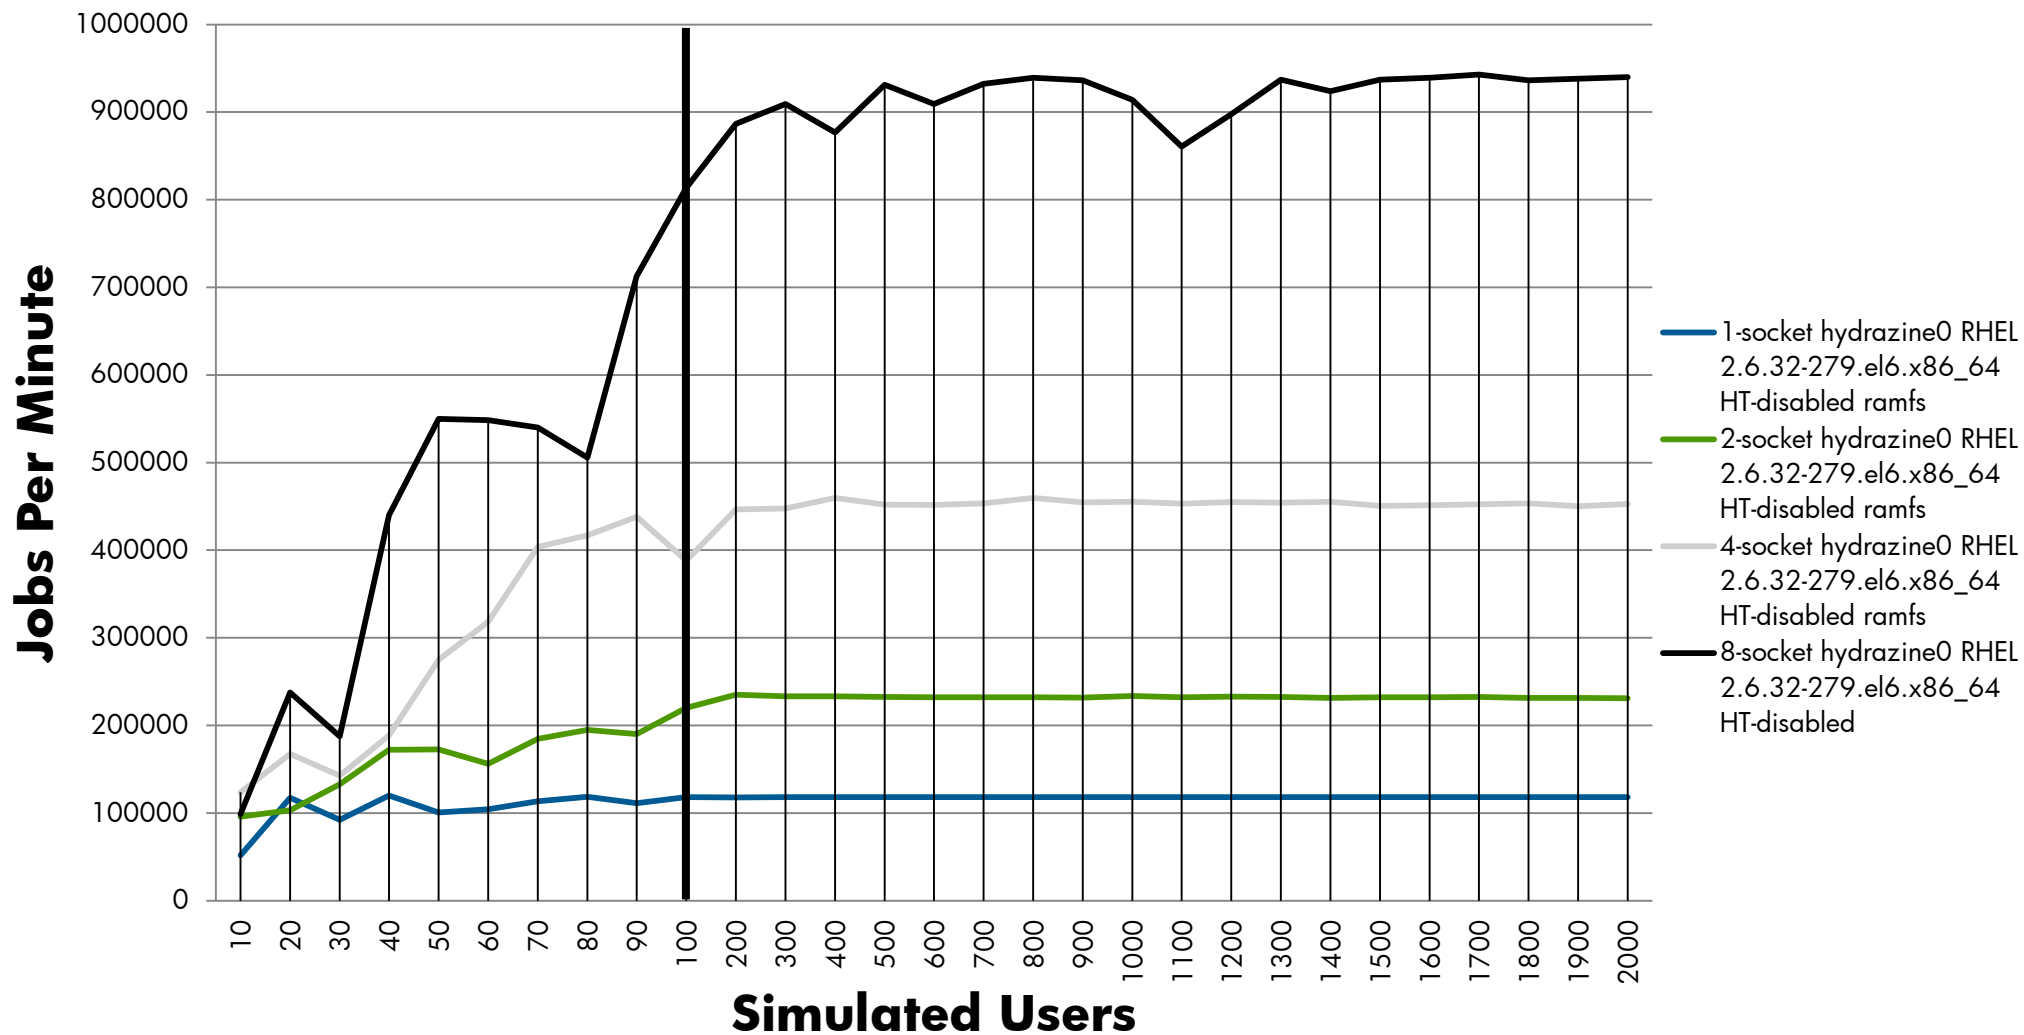


- Each successive run of `run_reaim` adds data to the end of the `<workload>.ssv` and `<workload>.csv` files
- This is done so that `<workload>.csv` can be incorporated into an Excel spreadsheet to graph the results
  - The same is done with the `<workload>.avg.csv` file
- Graphed Excel spreadsheet templates are located at
  - `/usr/local/share/reaim/results/aim_graph_template.xlsm`
  - `/usr/local/share/reaim/results/aim_graph_average_template.xlsm`
    - There is a worksheet tab for each of the 14 workloads in the Excel template
- The following slides show some sample graphs after incorporating `<workload>.csv` files to the template

# AIM7 – compute workload graph




# AIM7 - dbase workload graph





# AIM7 - fserver workload graph



# Graphing AIM-7 Output



- You should be really explicit and thorough about what you specify with the “-k <kernel>” option as this is what will distinguish your excel graph lines
  - In the future you may want to manually merge data from different runs
  - I’ve learned to specify things like HT-enabled/HT-disabled, ramfs/diskfs, power\_max/power\_min, etc in the -k <kernel> option rather than the -t <title> option because of this
  - Naming is important!

# Incorporating AIM-7 Output in Excel



- Instructions for Office 2007
  - Copy your `results` directory to a Windows machine (i.e., laptop or desktop)
  - Copy the `aim_graph_template.xlsm` template to the same directory
  - Open the Excel template: `aim_graph_template.xlsm`
    - If under the menu options (just above the opened file) you see:  
**Security Warning** some active content has been disabled Options...
      - Click the **Options...** box
      - A **"Security Alert"** dialog box will open
        - Select "Enable this content" in both the **"Macro"** and **"Data Connection"** sections
        - Click "OK"
        - Note: only need to do this when opening the file to update data, for viewing select "Help protect me"
  - Select the "UPDATE" worksheet at the bottom
  - Click on the "Press to Update Data" button
    - If any `.csv` file is missing the data in the tab for that workload will be empty – previous data is cleared
  - Save this file to a filename of your choice

# Incorporating AIM-7 Output in Excel



- In the presentation let's do a live example with:
  - `./sample`: Walk through a sample `.xslsm` file
  - `./add_to_excel`: Demo incorporating `.csv` data into `aim_graph_template.xslsm`
  - `cp ./copy_from/* add_to_excel/.` and incorporate the new data
  - `./average`: Demo incorporating `.avg.csv` data into `aim_graph_average_template.xslsm`
  - `./average`: Demo changing line color and style from average data that was put in `aim_graph_average_template.xslsm` and `aim_graph_template.xslsm`
  - `./average`: Demo incorporating the `.csv` data used to create the average data into `aim_graph_template.xslsm`
- Note: the `.xslsm` templates are sensitive to the format of the `.csv` files – don't change these output files

# How simple is it?



- To get the excel spreadsheet showing 1/2/4/8 socket DL980 scaling I wrote this script named `./my_per_socket_scr`:

```
./run_reaim -k "1-socket" -t "DL980" -c "numactl -m 0 -N 0"  
./run_reaim -k "2-socket" -t "DL980" -c "numactl -m 0,1 -N 0,1"  
./run_reaim -k "4-socket" -t "DL980" -c "numactl -m 0-3 -N 0-3"  
./run_reaim -k "8-socket" -t "DL980" # no need for numactl - run all
```

- Ran: `nohup ./my_per_socket_scr &`
- 12 hours later I copied the `./results` directory to my laptop
- Incorporated the `.csv` files into the Excel spreadsheet template
- All in all it took about 15-20 minutes of my time to write the script, start the run, copy the results and incorporate them into excel
- Obviously it will take longer when you are comparing different kernels that you will have to boot for each run!

# How simple is it? The Traditional Way



- To do the previous slide's socket scaling runs the "Traditional Way" would mean:
  - 14 AIM-7 workloads times 3 `reaim` commands per workload = 42 manual `reaim` commands
  - times 4 different socket configurations = 168 manual `reaim` commands
  - times 2 to also run with diskfs = 336 manual `reaim` commands (data not previously shown)
  - times 2 to also run with HT-enabled = 672 manual `reaim` commands (data not previously shown)
- Additionally, after each `reaim` command completes you have to manually save away the `multiuser.ss` and `reaim.csv` data files so the next run doesn't overwrite the data
- Finally you get to figure out how to pull all of that data together and manually create graphs to display all the results
- That's a lot of manual work by hand...

# How long does it take to run?



- All workloads:
  - 12-core DL380, RHEL 6.3, HT-disabled, ramfs: ~ 4 hrs
  - 12-core DL380, RHEL 6.3, HT-disabled, diskfs: ~ 4 hrs
  - 80-core DL980, RHEL 6.3, HT-disabled, ramfs: ~ 2.5 hrs
  - 80-core DL980, RHEL 6.3, HT-disabled, diskfs: ~ 4 to 5 hrs
- Individual workloads on an 80-core DL980, RHEL 6.3, HT-disabled, ramfs:

• all_ftime	~ 7 mins	• fserver	~ 13 mins
• alltests	~ 6 mins	• high_systime	~ 45 mins
• compute	~ 5 mins	• long	~ 7 mins
• custom	~ 8 mins	• new_dbase	~ 3 mins
• dbase	~ 3 mins	• new_fserver	~ 13 mins
• disk (no ramfs)	~ 7 mins	• shared	~ 5 mins
• five_sec	~ 14 mins	• short	~ 2 mins

# Backup Excel Slides





# Manually Incorporating Output in Excel



- Instructions for Office 2007
  - Open the Excel template: `aim_graph_template.xlsx`
  - Click on the “Data” menu tab
  - For each worksheet tab (e.g., “all\_untime”, “compute”, etc at the bottom)
    - Click on the worksheet tab you want
    - Select columns A-K, right-click, select “Clear Contents”
      - Click “Yes” on the pop-up dialog box
    - Click on cell A1
    - Click “From Text” on the Data menu tab – this will bring up an open dialog box
    - Navigate to the `.csv` file you want to incorporate and open the file
    - Click “Next” on first Import Wizard box (“Delimited” should be selected)
    - Click “Comma” on the second Import Wizard box, followed by “Next”
    - Click “Finish” on the third Import Wizard box (“General” should be selected)
    - In the Import Data box that appears, make sure it states “=\$A\$1”, click “OK”